

Verifiable Random Functions (Micali Rabin Vadhan, FOCS 1999): A Worked Summary with Constructions and Proof Ideas

Abstract

A verifiable random function (VRF) is a function keyed by a secret key that outputs, on any input x , a value y together with a non-interactive proof π that y is correct, such that y remains indistinguishable from random at any input for which no proof has been released. Micali Rabin Vadhan (FOCS 1999) formalize VRFs with a strong *unique provability* property and construct VRFs from an RSA root-extraction hardness assumption. The construction proceeds in three conceptual steps: (1) build a weaker primitive called a verifiable unpredictable function (VUF) from RSA; (2) lift VUFs to VRFs using a Goldreich Levin (GL) hardcore predicate idea; (3) extend fixed-length VRFs to all inputs $\{0, 1\}^*$ via a tree-based composition with prefix-free encoding.

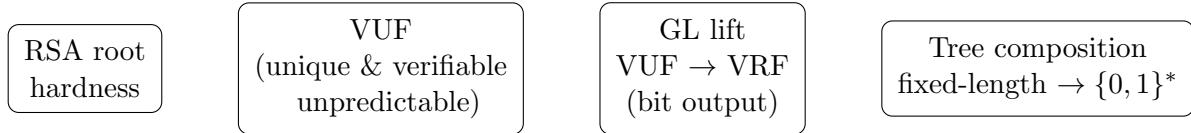
Contents

1 Roadmap (What is built and why)	2
2 Notation and background primitives	2
3 Primitives: VRFs and VUFs	3
3.1 VRFs (syntax and properties)	3
3.2 Residual pseudorandomness (the VRF security game)	3
3.3 VUFs: verifiable unpredictability	3
4 Hardness assumption used by the construction	4
4.1 RSA root-extraction with random large prime exponent	4
5 Construction 1: RSA-based VUF (value and proof)	4
5.1 High-level idea	4
5.2 Prime indexing: $x \mapsto p_x$	4
5.3 Algorithms	4
5.4 A compact correctness/uniqueness picture	5
5.5 Core properties and their proofs	5
5.6 Residual unpredictability: the reduction structure	6
6 Construction 2: Lifting VUFs to VRFs via Goldreich Levin	6
6.1 Construction (bit-valued VRF)	6
6.2 Why this yields pseudorandomness (proof idea)	6

7 Construction 3: Extending the input domain to $\{0,1\}^*$	7
7.1 Tree-based composition	7
7.2 Proof structure	7
7.3 Prefix-free encoding	8
7.4 Security idea (what must be shown)	8
8 Summary tables (what each step guarantees)	8
8.1 Primitive checklist	8
8.2 Construction and cost overview	9
9 End-to-end statement (what you obtain)	9
10 Practical reading notes (what to remember)	9

1 Roadmap (What is built and why)

The paper's approach can be read as a pipeline of reductions:



The remainder of this document:

- defines the primitives and security games precisely enough to reason about them,
- gives the concrete RSA-based VUF (including the proof/verification mechanism),
- explains the GL lift and why it yields pseudorandomness from unpredictability (with the key reduction idea),
- gives the domain-extension construction and its security intuition,
- summarizes the overall theorem and typical costs (proof size, verification work).

2 Notation and background primitives

Efficient algorithms and negligible functions. All algorithms are probabilistic polynomial-time (PPT) in the security parameter k . A function $\text{negl} : \mathbb{N} \rightarrow R_{\geq 0}$ is negligible if \forall polynomials p there exists k_0 such that for $k \geq k_0$, $\text{negl}(k) < 1/p(k)$.

Groups modulo an RSA modulus. Let $m = pq$ be an RSA modulus with distinct odd primes p, q . Let \mathbb{Z}_m denote integers modulo m and \mathbb{Z}_m^* its multiplicative group of units. Euler's totient is $\varphi(m) = (p-1)(q-1)$.

Lemma 1 (Permutation by exponentiation). *If $\gcd(e, \varphi(m)) = 1$, the map $\phi_e : \mathbb{Z}_m^* \rightarrow \mathbb{Z}_m^*$ defined by $\phi_e(z) = z^e \bmod m$ is a permutation.*

Proof. Because $\gcd(e, \varphi(m)) = 1$, there exists d with $ed \equiv 1 \pmod{\varphi(m)}$. Then for all $z \in \mathbb{Z}_m^*$, $(z^e)^d = z^{ed} \equiv z \pmod{m}$ by Euler's theorem, so ϕ_d is the inverse of ϕ_e . \square

Inner product over \mathbb{F}_2 . For $a, r \in \{0, 1\}^b$,

$$\langle a, r \rangle = \sum_{i=1}^b a_i r_i \pmod{2}.$$

3 Primitives: VRFs and VUFs

3.1 VRFs (syntax and properties)

A VRF is a triple (G, F, V) :

- $G(1^k) \rightarrow (\text{PK}, \text{SK})$.
- $F(\text{SK}, x) \rightarrow (y, \pi)$ where y is the function value and π is a proof.
- $V(\text{PK}, x, y, \pi) \in \{0, 1\}$ is the public verification algorithm.

Definition 1 (Correctness and complete provability). *For all x in the input domain, if $(\text{PK}, \text{SK}) \leftarrow G(1^k)$ and $(y, \pi) \leftarrow F(\text{SK}, x)$, then*

$$\Pr[V(\text{PK}, x, y, \pi) = 1] \geq 1 - \text{negl}(k).$$

Definition 2 (Unique provability). *For any public key PK (even adversarially chosen) and any input x , there do not exist two distinct values $y \neq y'$ with proofs π, π' such that $V(\text{PK}, x, y, \pi) = V(\text{PK}, x, y', \pi') = 1$, except with negligible probability.*

3.2 Residual pseudorandomness (the VRF security game)

Informally: even after seeing many (y, π) pairs at chosen inputs, the value at a fresh input looks random.

Definition 3 (Residual pseudorandomness game). *Let $\mathcal{O}_{\text{SK}}(x)$ return $(y, \pi) \leftarrow F(\text{SK}, x)$. Adversary \mathcal{A} gets PK and oracle access to \mathcal{O}_{SK} , outputs a fresh x^* (never queried), then receives either the real y^* or a uniform string of the same length; it may continue querying on inputs $\neq x^*$ and must guess which case it is in. The advantage is*

$$\text{Adv}_{\mathcal{A}}^{\text{vrf}}(k) = |\Pr[\mathcal{A} \text{ guesses correctly}] - \frac{1}{2}|.$$

A VRF is secure if this advantage is negligible for all PPT adversaries.

3.3 VUFs: verifiable unpredictability

The paper introduces a weaker primitive: a *verifiable unpredictable function* (VUF). The syntax, correctness, and unique provability are the same; the security goal is *unpredictability* rather than pseudorandomness.

Definition 4 (Residual unpredictability (VUF security)). *Given PK and oracle access to $\mathcal{O}_{\text{SK}}(x) = (y, \pi)$, no PPT adversary can output a fresh input x^* and a pair (y^*, π^*) such that $V(\text{PK}, x^*, y^*, \pi^*) = 1$, except with negligible probability.*

Remark 1. *A VUF is closely aligned with a unique signature scheme where the signature is the proof and the message is the input. The uniqueness condition matches unique provability.*

4 Hardness assumption used by the construction

4.1 RSA root-extraction with random large prime exponent

The core assumption is that extracting p -th roots modulo an RSA modulus is hard when p is a random large prime.

Definition 5 (RSA root hardness (informal)). *Let m be an RSA modulus of size $\approx k$ bits, let $x \leftarrow \mathbb{Z}_m^*$ be uniform, and let p be a random $(k+1)$ -bit prime (hence $p > m$). Given (m, x, p) , it is hard for any $s(k)$ -time adversary to find y such that*

$$y^p \equiv x \pmod{m}.$$

Why $p > m$ is useful for VRFs. Because $p > m > \varphi(m)$ and p is prime, we have $\gcd(p, \varphi(m)) = 1$, so the p -th root in \mathbb{Z}_m^* (if it exists) is *unique*. This uniqueness is exactly what the construction exploits to satisfy unique provability.

5 Construction 1: RSA-based VUF (value and proof)

5.1 High-level idea

Fix a public element $r \in \mathbb{Z}_m^*$. For each input x , associate a (publicly computable) prime exponent p_x . Define the function value to be the unique p_x -th root of r :

$$v_x := r^{1/p_x} \pmod{m}.$$

The proof is simply the witness v_x itself, and verification checks $v_x^{p_x} \equiv r \pmod{m}$.

5.2 Prime indexing: $x \mapsto p_x$

The construction needs a public map sending each x to a (nearly always distinct) prime p_x of about $(k+1)$ bits. The paper uses a *prime-sequence generator* (based on limited-independence polynomials and primality testing coins) to ensure that with overwhelming probability, all p_x are prime and distinct over the intended domain.

For this summary, treat the following as a black box:

Prime Indexer $\text{Prime}(x)$: a public deterministic algorithm (given some public seed) that outputs a $(k+1)$ -bit prime p_x for each input x , and outputs distinct primes for distinct x except with negligible probability.

5.3 Algorithms

Key generation.

$\mathsf{G}(1^k) :$ pick RSA modulus $m = pq$; pick $r \leftarrow \mathbb{Z}_m^*$; publish seed for $\text{Prime}(\cdot)$.

Public key:

$$\mathsf{PK} = (m, r, \text{seed}), \quad \mathsf{SK} = (\mathsf{PK}, \varphi(m)).$$

Evaluation (value + proof). On input x :

$$p_x := \text{Prime}(x), \quad d_x := p_x^{-1} \bmod \varphi(m), \quad v := r^{d_x} \bmod m, \quad \pi := v.$$

Output (v, π) (the proof is the same element).

Verification. On (x, v, π) (with $\pi = v$):

1. Compute $p_x := \text{Prime}(x)$; verify p_x is prime (with fresh randomness) and enforce $p_x > m$.
2. Check $v \in \mathbb{Z}_m^*$ and

$$v^{p_x} \equiv r \pmod{m}.$$
3. Accept iff all checks pass.

5.4 A compact correctness/uniqueness picture

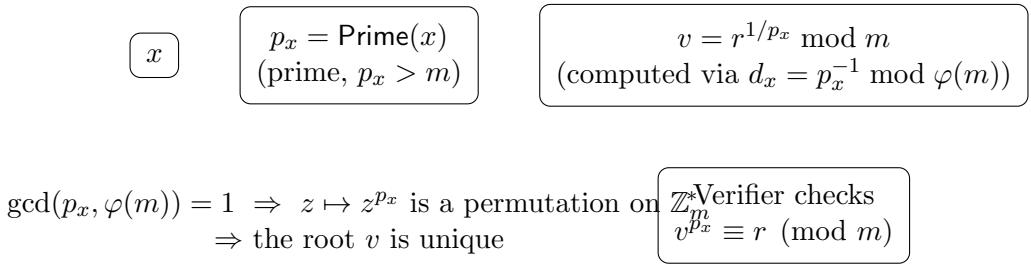


Figure 1: RSA-based VUF: the “proof” is the root v itself; verification is one modular exponentiation.

5.5 Core properties and their proofs

Lemma 2 (Completeness). *If $(v, \pi) \leftarrow \mathsf{F}(\mathsf{SK}, x)$ then $\mathsf{V}(\mathsf{PK}, x, v, \pi) = 1$ except with negligible probability.*

Proof. Let $p = p_x$ and $d = d_x$ with $pd \equiv 1 \pmod{\varphi(m)}$, so $pd = 1 + t\varphi(m)$ for some integer t . Then $v = r^d$ satisfies

$$v^p \equiv (r^d)^p = r^{pd} = r^{1+t\varphi(m)} \equiv r \cdot (r^{\varphi(m)})^t \equiv r \pmod{m},$$

since $r \in \mathbb{Z}_m^*$ implies $r^{\varphi(m)} \equiv 1 \pmod{m}$. □

Lemma 3 (Unique provability). *For any PK and input x , there is at most one $v \in \mathbb{Z}_m^*$ such that $v^{p_x} \equiv r \pmod{m}$ (except with negligible probability due to primality-test error).*

Proof. Verification enforces that p_x is prime and $p_x > m$, hence $p_x \nmid \varphi(m)$ because $\varphi(m) < m$. Thus $\gcd(p_x, \varphi(m)) = 1$ and exponentiation by p_x is a permutation on \mathbb{Z}_m^* . A permutation has at most one preimage for r , so at most one v can satisfy $v^{p_x} \equiv r$. □

5.6 Residual unpredictability: the reduction structure

The key security statement: producing a correct (v, π) for a fresh x^* is as hard as extracting a root with a fresh exponent. The paper's reduction uses two central ingredients:

1. **Programming one exponent:** choose a target x^* and construct the public seed so that $p_{x^*} = p$ where p is the RSA challenge exponent (while keeping the seed distribution close to honest).
2. **Algebraic simulation:** set $r := u^E \bmod m$ where u comes from the RSA challenge and $E = \prod_{x \neq x^*} p_x$. Then for any $x \neq x^*$ one can answer the oracle query by outputting

$$v_x := u^{E/p_x} \bmod m,$$

since $(v_x)^{p_x} = u^E = r$.

If the adversary outputs a valid v^* for x^* , then $(v^*)^p = r = u^E$. Using Bézout coefficients α, β with $\alpha E + \beta p = 1$, one extracts $u^{1/p}$ as

$$u^{1/p} \equiv (v^*)^\beta \cdot u^\alpha \pmod{m}.$$

Theorem 1 (RSA-based VUF security (informal)). *Assuming RSA root-extraction is hard for random large primes, the above construction is a secure VUF: any PPT adversary that forges a valid value at a fresh input with non-negligible probability yields an RSA-root inverter with non-negligible probability (up to the standard “guess the challenge input” factor when selecting x^*).*

6 Construction 2: Lifting VUFs to VRFs via Goldreich Levin

6.1 Construction (bit-valued VRF)

Let the VUF output be encoded as a b -bit string $v(x) \in \{0, 1\}^b$. Sample a public random $r \in \{0, 1\}^b$ and define the derived output bit:

$$y(x) := \langle v(x), r \rangle \bmod 2.$$

Evaluation returns a proof that reveals $v(x)$ and proves it is correct:

$$\mathsf{F}'(\mathsf{SK}, x) : (v, \pi) \leftarrow \mathsf{F}(\mathsf{SK}, x), \quad y = \langle v, r \rangle, \quad \pi' = (v, \pi).$$

Verification checks both:

$$\mathsf{V}'(\mathsf{PK}', x, y, \pi') = 1 \iff (\mathsf{V}(\mathsf{PK}, x, v, \pi) = 1) \wedge (y = \langle v, r \rangle),$$

where $\mathsf{PK}' = (\mathsf{PK}, r)$ and $\pi' = (v, \pi)$.

6.2 Why this yields pseudorandomness (proof idea)

Goldreich Levin (GL) says: if you can predict $\langle w, r \rangle$ for random r with noticeable advantage, you can reconstruct w with non-negligible probability (using $\text{poly}(k)$ queries to the predictor).

Here, the role of w is played by $v(x^*)$ at a fresh input x^* . So, if an adversary distinguishes $y(x^*)$ from uniform, one can convert it into a predictor for $\langle v(x^*), r \rangle$, then apply GL reconstruction to recover $v(x^*)$ itself. But recovering $v(x^*)$ (together with its proof) breaks the VUF unpredictability.

Assume distinguisher \mathcal{D}
distinguishes $y(x^*) = \langle v(x^*), r \rangle$ from uniform.

Build predictor \mathcal{P}
that guesses $\langle v(x^*), r \rangle$
with noticeable bias.

Recovered $v(x^*)$ yields
a correct fresh VUF value
(breaks unpredictability).

Goldreich Levin
reconstructs $v(x^*)$
from biased inner-product oracle.

Figure 2: Security flow for the GL lift: distinguisher \Rightarrow predictor \Rightarrow GL reconstruction \Rightarrow VUF break.

The technical nuisance addressed in the paper: the adversary chooses x^* adaptively, while GL wants a *fixed* $w = v(x^*)$. The reduction handles this by choosing a random x^* and hoping it matches the adversary’s exam point (the familiar $2^{-|x|}$ loss), motivating an initial “short-input” regime; the next construction removes this restriction.

Theorem 2 (VUF \Rightarrow VRF (informal)). *If $(\mathbf{G}, \mathbf{F}, \mathbf{V})$ is a secure VUF with unique provability, then the GL-derived $(\mathbf{G}', \mathbf{F}', \mathbf{V}')$ is a (bit-valued) VRF for an appropriate parameter regime (with polynomial overhead and the standard challenge-point guessing loss).*

7 Construction 3: Extending the input domain to $\{0, 1\}^*$

7.1 Tree-based composition

Assume a base VRF (or verifiable pseudorandom predicate) that can be applied iteratively. The idea is to label a binary tree so that each node label deterministically defines its children via the base VRF.

Conceptually:

- Root has label $L(\epsilon)$.
- For a node with label $L(w)$, define

$$L(w0) := f(L(w), 0), \quad L(w1) := f(L(w), 1),$$

where f is a base VRF/predicate-to-string variant.

- For $x = b_1 \cdots b_t$, the output is $L(x)$, the label at the end of the path.

7.2 Proof structure

A proof for $x = b_1 \cdots b_t$ contains the intermediate labels and per-edge proofs:

$$(L(b_1), \pi_1), (L(b_1 b_2), \pi_2), \dots, (L(b_1 \cdots b_t), \pi_t),$$

where π_i proves that $L(b_1 \cdots b_i)$ was computed correctly from $L(b_1 \cdots b_{i-1})$ and bit b_i .

7.3 Prefix-free encoding

To prevent revealing values on prefixes of future challenges, the input x is first mapped to a prefix-free encoding \hat{x} (so no \hat{x} is a prefix of \hat{x}' for $x \neq x'$). The tree walk is done on \hat{x} instead of x .

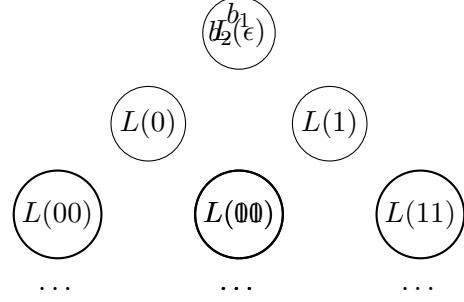


Figure 3: Domain extension: compute $L(\hat{x})$ by walking the tree. Proofs certify each edge transition.

7.4 Security idea (what must be shown)

Fix an adversary that sees many path proofs and then challenges on a fresh \hat{x} . Two complementary arguments are used:

- **Conditioned-on-no-collisions:** if no label repeats among nodes revealed so far, then the challenge label is distributed like a fresh base-VRF output at an unseen point, hence pseudorandom.
- **Collisions are unlikely (or exploitable):** if label repetitions happen with noticeable probability, one can leverage the first collision to predict a new label and contradict the base VRF security (label lengths are chosen to push collision probability down, and any non-negligible collision probability yields a distinguisher/predictor).

Theorem 3 (Fixed-length \Rightarrow unrestricted-length VRF (informal)). *Given a secure fixed-length VRF with sufficiently long labels (output length), the tree construction (with prefix-free encoding) yields a secure VRF on $\{0, 1\}^*$, with proof size and verification time linear in $|\hat{x}|$.*

8 Summary tables (what each step guarantees)

8.1 Primitive checklist

Primitive	Verifiable?	Unique proof?	Security goal
VUF	yes	yes	unpredictability at fresh points
VRF	yes	yes	pseudorandomness at fresh points
Tree-extended VRF	yes	yes	pseudorandomness on $\{0, 1\}^*$

Table 1: Conceptual distinction: VUF vs. VRF.

8.2 Construction and cost overview

Step	What is built	Main cost
$\text{RSA} \Rightarrow \text{VUF}$	$v_x = r^{1/p_x} \bmod m$ with proof $\pi = v_x$	One exponentiation for verify; indexer computes p_x
GL lift	Bit output $y = \langle v, r \rangle$ with proof revealing v	Proof includes v (and VUF proof), plus GL reduction overhead
Tree extension	VRF on $\{0, 1\}^*$ via path labels and proofs	Proof size $\Theta(\hat{x})$; verify $\Theta(\hat{x})$

Table 2: Each reduction adds structure (and typically proof length) while preserving verifiability and uniqueness.

9 End-to-end statement (what you obtain)

Theorem 4 (End-to-end outcome (informal)). *Assuming RSA root-extraction is hard for random large prime exponents, there exists a VRF family on inputs $\{0, 1\}^*$ that outputs at least one pseudorandom bit together with a publicly verifiable, uniquely valid proof of correctness. The construction is explicit: RSA-based VUF \Rightarrow GL-derived verifiable pseudorandom predicate \Rightarrow tree-based domain extension.*

10 Practical reading notes (what to remember)

- **The proof is not a separate object in the RSA VUF:** the witness v_x is the proof.
- **Uniqueness is structural, not heuristic:** enforcing p_x prime and $p_x > m$ makes exponentiation a permutation.
- **Security reductions have two recurring motifs:** (i) *program a special input* and pay a “guess the challenge” loss, (ii) *algebraic simulation* that answers all other queries consistently without knowing secret structure.
- **Domain extension trades interaction for proof length:** proofs grow linearly with input length.